

COMPANY · 7 MIN READ

The Linguini business model


Free for a daily photo, Plus when one a day isn't enough. How we priced Linguini around learners, competitors and the AI cost of every lesson.

Free gives every learner **one photo lesson and one journal page a day**. Plus is **\$49.99 a year** for anyone who wants more. Each photo lesson costs us about **2¢ in AI**, so the free tier pays for itself once enough people upgrade: **about 4% of learners** on the models we chose, or under 2% once cheaper models pass our tests.

The model at a glance

The free plan limits the one thing that costs us money: AI reading your own photos. The learning stays free.

■ How we deliver ■ What learners get ■ Who we serve ■ Money



Linguini's business model canvas. The customer segments are example learner types, not real users.

Who it is for

Our learners are people returning to a language they half know, beginners who want a few minutes a day, and students and travellers. None of them want to sit through a course, and most will never pay for one: across education apps, only 2.3% of downloads turn into a paid subscription.

So the free plan has to be good enough to become a habit on its own, and the reason to upgrade has to show up at the moment someone wants more. For us that moment is easy to spot. It is the day you take more photos than one lesson covers.

Plans

	Free \$0 forever	Plus \$49.99 a year · or \$7.99/mo	Founding \$34.99 a year · first 300
Photo lessons	1 /day	10 /day	10 /day
Photos per journal page	1	10	10
Curated scenes	∞	∞	∞
Pronunciation feedback	—	✓	✓
Review mode	—	✓	✓
Competitive I-Spy beta	—	—	✓
Vote on the roadmap	—	—	✓

Curated scenes are analysed once in advance, so replaying one costs us about a sixth of a cent and free learners can play them all day. Plus adds the things that cost us money: more photo lessons, more photos per journal page and feedback on your speaking.


Founding is Plus for our first 300 members, at \$34.99 a year for as long as they stay. We give up some margin on them. In return we get early payers who tell us what to build, and in the first year that is worth more than the margin.

How we set the price

There are three usual ways to price something: from what it costs you, from what competitors charge, and from what it is worth to the customer. We checked \$49.99 all three ways.

What a year of Linguini should cost

Annual prices, US dollars



Cost sets the floor.

A typical Plus learner costs us about \$15 a year in AI, storage and card fees. Anything above that can also carry some free learners.

Competitors set the range.

Photo-flashcard apps charge \$20–30 a year. Duolingo, Speak and Babbel charge \$84–168. We do more than a flashcard app and less than a full course.

Value sets the ceiling.

An online tutor costs \$15–60 an hour. A year of daily practice should cost less than one or two lessons.

So Plus is \$49.99 a year.

That is close to the \$44.99 education-app median and below every course app. Our first 300 members lock in \$34.99.


Against the competition

App	Price a year	Own photos	Journal
Linguini Plus Free: 1 photo lesson a day	\$49.99	☒	☒
CapWords Premium Free: A few captures	\$19.99–29.99	☒	☐
Speak Premium Free: Short taster	\$83.99	☐	☐
Duolingo Super Free: About 2–3 lessons a day	\$84–96	☐	☐
Babbel 12 months Free: Short taster	\$107.64	☐	☐

US list prices checked September 2026; they vary by country and with discounts. CapWords turns photos into flashcards, without games or a journal.

CapWords is cheaper, but it stops at flashcards. The course apps do more teaching and charge nearly twice as much. We sit between them, so our price does too. From Duolingo we took the shape of the free tier: keep it genuinely useful, limit how much you can do, and save the expensive AI for paid plans.

What a lesson costs us



Which model runs each call

Every call goes through OpenRouter, with one key, and each call has its own setting. We picked each model by running the candidates through the app's own prompts and checks: 14 labelled photos, 17 translations, 6 scenes and 22 learner guesses. No provider won everywhere, so we use three. That has a second benefit: if one provider goes down, one feature falls back to its alternative and the rest of the app carries on.

Call	Chosen, via OpenRouter	Main alternatives	Why
Scene analysis	Claude Haiku 4.5 (Anthropic)	GPT-4o, GPT-4.1-mini	Best quality, found every object we labelled, and the only model the app accepted every time.
Translation	GPT-4o-mini (OpenAI)	Mistral Small, Claude Haiku 4.5	Nearly as good as Mistral at the same price, without the throttling Mistral hit.
Learning tasks	GPT-5.4-mini (OpenAI)	GPT-4.1-mini, Gemini 3.5 Flash-Lite	The only model that reliably fits our lesson format. Also our most expensive call.
I-Spy clues	Gemini 3.1 Flash-Lite (Google)	GPT-4o-mini, Claude Haiku 4.5	The best clues, in perfect target language, and it never gave the answer away.
I-Spy guesses	GPT-4.1-mini (OpenAI)	Gemini 3.5 Flash-Lite, GPT-4o-mini	The most accurate at working out which object a learner meant.

Chosen after our model comparison. Temperature is 0 for every call, because raising it made every call worse.

What a learner costs

A lesson costs about 2¢, so what a learner costs us depends on how much they use Linguini. The table shows a month for different kinds of learner.

Learner	Photo lessons a month	Chosen models	Cheapest alternatives
Free, light	2	\$0.04	\$0.02
Free, casual	8	\$0.17	\$0.07
Free, every day	30	\$0.64	\$0.27
Plus, typical	30, with speaking	\$0.74	\$0.37
Plus, heavy	90	\$2.17	\$1.06
Plus, 10 a day, every day	300	\$7.13	\$3.47

AI and storage per learner per month, at OpenRouter prices. The cheapest alternatives are not ready yet: they score lower in our tests.


After Stripe’s fees, a Plus learner brings in about \$5.24 a month. The typical payer is comfortably profitable, and so is the heavy one. We only lose money on someone who uses all ten lessons every day. That is why Plus stops at ten a day instead of promising “unlimited”: unlimited means losing money on your most enthusiastic learners.

Hosting is fixed at \$52–70 a month (Vercel, Supabase and Render), and at every size we modelled it is smaller than the AI bill. So our costs grow with use, not with signups, and the daily limits are what keep them in check. Stripe also keeps 3.4% + 50¢ of each payment, which is why we push the annual plan.

Why freemium

Ordinary software costs almost nothing to run for one more user. AI doesn’t: every lesson is a bill. Many AI companies charge less than cost and bet that models will get cheaper. We can’t make that bet. A cheaper model only helps if it still makes a good lesson, and in our tests the cheapest ones did not, yet. The pricing has to work at the costs we actually face.

We judged each option on two questions: does it cover the AI cost, and will learners like it?




- **Freemium with daily limits (chosen).** Free learners get one photo lesson a day, which caps what they cost. Everyone can try it first.
- **Usage credits.** Matches our costs exactly, but counting credits breaks a daily habit.
- **Flat subscription.** No way to try before paying.
- **Lifetime deal.** Popular, but AI costs keep coming after the one payment.
- **Ads.** Earns little at our size, and ads would sit next to personal photos.

Does the free tier pay for itself?

Every payer also carries the free learners: about 42 of them when 2.3% pay. On the models we chose, that only pays off once about 4% of learners pay. The cheapest alternatives would break even at under 2%, but they make worse lessons today, so we launch on the chosen models and work on bringing the cost down.

With **5,000** learners and **2.3%** paying, Linguini makes **-\$447** a month on the chosen models.

On the cheapest alternatives, once they pass our evaluation, it would be +\$126.



Before a wide launch we will do three things. First, fix the defect in our lesson format that makes most models fail it, so the lesson generator can move to GPT-4.1-mini at less than half the cost. Second, run the cheapest alternatives through the same evaluation and switch only the calls that pass. Third, enforce the daily limits on the server and track what every lesson costs.

Where the money comes from

Plus is the business. Everything else either brings in money early or comes later, once we know learners stay.

Stream	Price	When	Why
Plus	\$49.99 a year or \$7.99 a month	At launch	The core. Sold on the web, so no app store takes a cut.
Founding Plus	\$34.99 a year, first 300 members	Launch, until 300 members or 31 Dec 2026	Early revenue from people who will tell us what to build. Each nets about \$2.74 a month against \$1.03 in AI.
Family plan	About \$79.99 a year for up to 4	Once retention is proven	Every member uses AI, so we price per seat instead of copying Duolingo's \$119.99 for 6.
Classroom licences	Not set	Year two	A teacher sets a daily scene and the class journals about it.
Printed journal	Needs a print quote	Once PDF export ships	A one-off purchase that fits the idea of keeping the day.

What success looks like, stage by stage

Revenue is the wrong measure for a product nobody comes back to. In the beta, one learner who returns tells us more than one who pays once and leaves. So the number we watch changes as the question changes: first whether people come back, then whether they pay, then whether enough of them pay to carry everyone else.

Stage	The question	What we measure	Money coming in
Now, in beta	Does it stick?	New learners who take a second photo lesson within a week	None. The beta costs us about \$75 a month, which we treat as research.
Launch year	Will people pay?	300 Founding Plus members	Plus and Founding Plus
After that	Does it scale?	4% of monthly learners paying	Plus, then family plans, classrooms and printed journals

Sources

[OpenRouter model prices](#) · [Supabase](#) · [Vercel](#) · [Stripe](#) · [Duolingo Q2 2026 letter](#) · [Duolingo prices](#) · [Babbel prices](#) · [Speak prices](#) · [CapWords](#) · [RevenueCat SOSA 2026](#) · [italki tutor costs](#). Token counts and quality scores come from our model comparison, which ran every candidate through the app's own prompts, schemas and validators.